

Westpac Sales Scraper

How to use the dashboard — in plain English, for the people who run it every day.

What this tool does: it finds businesses (contractors and material distributors), cleans the list, adds contact details, and gives you a ready-to-use sales sheet.

Who this guide is for: end users — the people clicking buttons in the dashboard. No engineering knowledge needed.

Florida edition **EXISTING** · Tennessee + Vendors **NEW**

What's inside

1. **The big picture** — what the tool does, in one screen
2. **Three kinds of runs** — Florida contractors, Tennessee contractors, Tennessee vendors
3. **The Dashboard** — every button and setting explained (click this → this happens)
4. **How a run works** — the 5 pipeline stages
5. **Florida contractors** OLD — unchanged
6. **Tennessee contractors** UPDATED logic — vendors first, then geography
7. **Tennessee vendors** NEW
8. **The menu** — what each page is for
9. **Tags** — the labels on every record, and what data each one carries
10. **Click a record** — the detail panel (Tags + where the data came from)
11. **Data separation** — raw → staged → final, and why some rows disappear
12. **Requirements** → **where they live** — the client's brief mapped to the UI
13. **Quick answers** — common questions

1 · The big picture

The scraper turns "go find me sales leads" into a clean spreadsheet, automatically.

You tell it **what** to find (contractors or vendors) and **where** (Florida or Tennessee). It then searches Google Maps, removes duplicates and businesses you don't want (like lumber yards), adds phone/email/license details, and saves a final list you can export.



THE GOLDEN RULE OF THIS UPGRADE

Everything that worked before for **Florida contractors** still works exactly the same. Tennessee and Vendors were **added on top** — nothing old was changed or removed.

Two new things you'll notice

- A **"Vendors"** section in the menu, separate from Contractors. Vendors are the material distributors you sell to (drywall/gypsum suppliers).
- **Tags on every record** — small labels like `state: TN` `Tier 1` `contractor` that tell you exactly what a row is and let you filter the list.

2 · Three kinds of runs

On the Dashboard you choose two things — **What** (Contractors / Vendors) and **Where** (Florida / Tennessee). That gives three valid combinations:

Run	What it finds	Status
Florida · Contractors	Drywall/construction contractors across Florida cities — the original product.	OLD · unchanged
Tennessee · Contractors	Contractors in Tennessee, searched around your vendor locations (new geography logic).	UPDATED
Tennessee · Vendors	Drywall-material distributors (GMS, L&W Supply, etc.) around Tennessee city centers.	NEW

IMPORTANT ORDER FOR TENNESSEE

For Tennessee, **run Vendors first, then Contractors**. The vendor locations become the anchor points that the contractor search circles around. (Full explanation in sections 6 & 7.)

WHY IS "FLORIDA · VENDORS" MISSING?

Vendor scraping is a **Tennessee-only** feature by design. If you switch to Vendors mode, Florida is automatically disabled — that is expected, not a bug.

3 · The Dashboard — every setting

This is the screen where you start runs and tune settings. Here is each part, top to bottom, with what happens when you use it.

Max records (testing limit)

Dashboard · Max records to discover

A number box. It caps how many businesses the very first search step collects.

Click/type a small number (e.g. 10) → quick test run, finishes fast, low cost.
Leave it high / blank → a full production run. Use small numbers while testing, then raise it for the real run.

Data source toggles (Apify · BBB · Apollo)

Dashboard · Data sources

Three on/off switches:

- **Apify — Google Maps (Discovery)**: the main "find businesses" engine.
- **Apify — BBB (Enrichment)**: adds Better Business Bureau ratings/accreditation.
- **Apollo (Enrichment)**: looks up business email addresses.

Turn one off → that step is skipped on the next run (e.g. turn off Apollo to skip email lookup and save credits).
All on → the most complete data.

Tennessee search radii (miles) TN only

Dashboard · Tennessee search radii

Two number boxes. They control how wide the Tennessee searches reach.

Setting	Default	What it controls
Vendor radius	20 miles	When scraping vendors , how far around each Tennessee city center to look.
Contractor radius	50 miles	When scraping TN contractors , how far around each vendor location to look.

These are two different radii because they're used by two different searches. They only affect Tennessee — Florida ignores them completely. Changes apply to the **next** run.

Statewide TN license verify (optional) TN only

Dashboard · Statewide TN license verify

A checkbox. Tennessee has no single free statewide license file, so this is off by default.

Leave unchecked (default) → fast TN runs using the free Nashville license data.

Check it → also cross-checks the state verify site for extra license matches. Slower. Florida is unaffected either way.

Save settings button

Saves the radii / verify / source choices. The button only lights up when you've actually changed something, so you always know if there are unsaved edits.

Run scope — What & Where

Dashboard › Run scope

Two rows of buttons:

- **What:** **Contractors** **Vendors**
- **Where:** **Florida** **Tennessee**

Click "Vendors" → Where automatically jumps to Tennessee and Florida greys out (vendors are TN-only).

Click "Contractors" → you can pick either Florida or Tennessee.

HELPER NOTE YOU'LL SEE

When you pick **Tennessee + Contractors**, the dashboard shows a reminder: the run searches the **contractor radius around your vendor accounts**. If no vendors exist yet, it falls back to searching TN cities directly (Florida-style).

The Start button

The big button reads exactly what you're about to do, e.g. **"Start TN Vendor Scrape"** or **"Start FL Contractor Scrape"**. Click it and the run begins; you can watch progress on the **Pipeline Stages** and **Logs** pages.

4 · How a run works — the 5 stages

Every run, no matter the type, moves through the same five stages. The **Pipeline Stages** page shows them live so you can see where the data is.



#	Stage	What happens
1	Discovery	Searches Google Maps for the businesses in scope and collects the raw list.
2	Dedupe	The same business found by two searches is merged into one.
3	Classify	Labels each business (its tags) and flags lumber/out-of-area ones for removal.
4	Cap	Applies your "max records" limit so the list doesn't get larger than asked.
5	Enrich + Save	Adds BBB/license/Apollo details and saves the finished records (contractors → Contractors, vendors → Vendors).

On the Pipeline Stages page these appear as "1 · Discovery", "2 · Dedupe", ... so you can click any stage and see exactly which records made it through.

5 · Florida contractors

OLD — unchanged

This is the original product. It behaves exactly as it always did.

When you run **Florida · Contractors**:

1. The tool goes through the **Florida cities** list and their ZIP codes.
2. It searches Google Maps in each of those areas for drywall/construction contractors.
3. It removes duplicates, drops lumber businesses, adds details, and saves them to the **Contractors** page.

NOTHING HERE CHANGED

Same cities, same keywords, same results as before the upgrade. Florida does **not** use vendors, radii, or the TN license verify. Old records may show no **state** tag — that's normal; they are treated as Florida.

6 · Tennessee contractors

UPDATED logic

Tennessee contractors are found using a new "vendors first" geography — not the Florida city method.

The new idea: search around your vendors

In Tennessee, your best contractor leads are the ones **near the distributors that sell materials**. So instead of sweeping whole cities, the tool draws a circle (the **contractor radius**, default 50 miles) around **each vendor location you already scraped**, and searches inside those circles.



THIS IS WHY ORDER MATTERS

The vendor-geography logic **only runs if vendors already exist**. So the correct sequence is: **(1) run Tennessee Vendors → (2) then run Tennessee Contractors**.

What if you skip the vendor run?

If there are **no vendors yet**, the tool doesn't fail — it **falls back to the Florida-style method**: it searches the Tennessee cities and their ZIPs directly, just like Florida does. You still get contractors; they just aren't anchored to vendor locations.

HOW TO TELL WHICH ONE RAN

The **Logs** page tells you. If you see a line about the *Florida-style / city fallback* during a **TN contractor** run, it means no vendors were available, so the vendor-geography logic was skipped. To get the intended

behaviour, run a vendor scrape first and try again.

Situation	What the tool does
Vendors already scraped for TN	Vendor-geography: searches the contractor radius around each vendor.  intended
No vendors yet	Fallback: searches TN cities/ZIPs directly (Florida-style). Works, but not anchored to vendors.

7 · Tennessee vendors NEW

Vendors are the material distributors you sell to — and they live in their own section.

What a "vendor" is here

A vendor = a **drywall / gypsum / building-material distributor** (e.g. GMS, L&W Supply). These are different from contractors, so they are stored on a **separate Vendors page**, not mixed into Contractors.

How the vendor run works

1. It takes the **Tennessee city centers** and draws the **vendor radius** (default 20 miles) around each.
2. It searches Google Maps inside those circles for material distributors.
3. It keeps only the real distributors and **drops contractors that look like suppliers** (e.g. "S & L Carpentry" is a contractor, not a vendor — it's removed).
4. It rolls related branches up under one **network name** (many local GMS branches → one "GMS"), and flags large chains as **big-box**.
5. The clean distributor list is saved to the **Vendors** page.

WHY VENDORS ARE SEPARATE

Contractors and vendors are two different audiences. Keeping them in separate pages/tabs means a vendor never accidentally shows up in your contractor sales sheet, and the vendor list is also what powers the Tennessee contractor geography (section 6).

What the Vendors page shows

A table just like Contractors — distributor name, **Network**, **Type** (or a **big-box** badge), address, city, **Tier**, ZIP, phone, email, website — with a filter box on every column and an **Export CSV** button.

8 · The menu — what each page is for

Menu item	What you do there
Dashboard	Pick what/where, tune settings, and start a run (section 3).
Cities	The list of cities the tool searches, with their tier, county and search radius. Tennessee cities are seeded automatically.
Vendor Aliases	The rules that roll branch names up into one network (e.g. all "L & W" spellings → "L&W Supply").
Keywords	The search terms used to find contractors (seeded automatically on a fresh sheet).
Contractors	Your final contractor results — filter, inspect, and export. (This was called "Results" before.)

Vendors NEW	Your final distributor results — separate from contractors.
Pipeline Stages NEW	See a run move through the 5 stages and inspect what survived each one.
Logs	Live text log of the current/last run — this is where the "which logic ran" messages appear.
History	Past runs, when they happened, and how many records each produced.

Note: the old "Dealer Accounts" page was removed — "dealer", "vendor" and "distributor" all mean the same thing here, and they now live under **Vendors**.

9 · Tags — the labels on every record

Tags are how the tool "separates by labels instead of making copies." Every contractor and vendor carries them, and you filter the list using them.

You'll see these as small coloured badges. Here is what each one means and what kind of data it holds:

Tag	What it means	Data you'll see
record_type	Is this row a contractor or a vendor?	<code>contractor</code> · <code>vendor</code> (blank = old Florida contractor)
state	Which state the business is in.	<code>FL</code> · <code>TN</code>
county	The county within the state.	e.g. <code>Davidson</code> , <code>Shelby</code>
city	The city.	e.g. <code>Nashville</code> , <code>Memphis</code>
city_tier	How important the city is (1 = top priority market).	<code>1</code> · <code>2</code>
zip_code	The postal ZIP of the business.	e.g. <code>37203</code>
source	Where this record was found.	<code>google_business</code> · <code>vendor_seed</code> · <code>bbb</code> · <code>license</code>
canonical_network	Vendors: the parent network branches roll up into.	e.g. <code>GMS</code> , <code>L&W Supply</code>
vendor_type	Vendors: what kind of distributor.	<code>specialty_distributor</code> · <code>big_box_retailer</code> · <code>independent</code>
is_big_box	Vendors: is this a large chain store?	<code>big-box</code> badge when true
out_of_territory	Outside the area you sell to → kept for audit, hidden from exports.	true / false
excluded_reason	Why a row was removed (mostly lumber).	e.g. <code>lumber:category:lumberyard</code>
enrichment_status	How far detail-adding got for this row.	e.g. <code>enriched</code> / <code>partial</code>
canonical_entity_id	The tool's internal "same business" ID — links a record to its raw sources.	an ID string

HOW TO USE TAGS TO FILTER

On the Contractors and Vendors pages, every column has a **filter box**. Type into the **state**, **city**, **Tier**, **Network** or **Type** box to instantly narrow the list. The tags are exactly those filterable fields, shown as badges so you can see them at a glance.

10 · Click a record — the detail panel

Click any contractor or vendor row and a panel slides in from the side, explaining exactly what that record is and where its data came from.

Contractors / Vendors › click a row → detail drawer

TAGS SECTION

Shows every label on this record as badges — **contractor** **TN** **Nashville** **Tier 1** **37203** **GMS** — so you instantly understand **what** the business is and **how** it was categorised.

SOURCES (RAW LAYER) SECTION

Lists each original source row that fed this record — e.g. the Google Maps result, the BBB listing, the license match. This is the **provenance**: it answers "where did this data come from?" Every source is linked to the record by its `canonical_entity_id`.

CLASSIFICATION SECTION (CONTRACTORS ONLY)

For contractors, an extra "why was this included" audit trail. (Vendors don't show this — it's contractor-specific.)

So the flow is: **see a row** → **click it** → **read its Tags (what it is) and its Sources (where it came from)**. That's the whole "data separation" idea made visible — one clean record on top, its raw source rows underneath.

11 · Data separation — raw → staged → final

The tool keeps three layers of data so nothing is lost and your final list stays clean.



Layer	What it is	Where you see it
A · Raw sources	Every individual result, kept exactly as found, one row per source (Google, BBB, license...).	The Sources section in a record's detail panel (section 10).
B · Staged	A snapshot of the records after each of the 5 stages.	The Pipeline Stages page.
C · Final	The clean, deduped, enriched list — your deliverable.	The Contractors and Vendors pages and the CSV export.

Why some rows "disappear" from exports

The final export is the **clean sales list**, so by default it hides:

- excluded_reason** rows — mostly **lumber** businesses, which you don't sell to.
- out_of_territory** rows — businesses outside your selling area.

They aren't deleted — they're tagged and set aside. You can still include them for auditing by exporting with the **"include excluded"** option.

ABOUT LUMBER

"Lumber" businesses (lumber yards, sawmills, etc.) are intentionally excluded because they aren't drywall-material targets. The tool is careful here — for example **"plumber"** is *not* treated as lumber. Removed rows always show a clear **excluded_reason** tag so you know why.

12 · Requirements → where they live

This connects the client's brief to the running product — so anyone can see each requirement and exactly where to find it in the dashboard.

SOURCE DOCUMENT

The requirements come from **"Scraper Platform Upgrade — Developer Implementation Checklist"** (the engineering brief). It groups the work into five **workstreams (A–E)**. The brief's golden rule: *"This is*

an additive release. Existing Florida behavior must keep working unchanged.” — which is exactly how the product behaves.

Legend: **DONE** built & in the UI · the “Where to see it” column is the page/section in this guide.

Workstream A — Geography & territory (add Tennessee)

Requirement (from the brief)	Where to see it in the dashboard	
Add Tennessee alongside Florida (Florida untouched)	Dashboard › Run scope › Where — Florida / Tennessee buttons (§3)	DONE
Move locations into editable config / reference tables	Cities page — editable city/tier/county/radius list (§8)	DONE
Region include/exclude list, editable at run time	Cities / exclusions — Memphis metro is excluded automatically	DONE
Hard exclusion: Memphis & metro never scraped	Applied during a TN run; excluded rows carry out_of_territory (§9, §11)	DONE
Two configurable radii (vendor 20mi, contractor 50mi)	Dashboard › Tennessee search radii (§3)	DONE
Geocoding step (addresses → coordinates)	Runs automatically inside a TN run (no button needed)	DONE
City tiers: Tier 1 priority cities, Tier 2 by population	Cities page tier column; city_tier tag on every record (§9)	DONE
Tier as both run-order and a tag on every record	Tier 1 runs first; tag visible in Contractors/Vendors + detail box (§9, §10)	DONE

Workstream B — Tennessee contractor scraper

Requirement (from the brief)	Where to see it in the dashboard	
Do not change the proven keyword/search logic	Same engine as Florida; TN only feeds it different areas (§5, §6)	DONE
Feed it the TN in-territory ZIP set from Workstream A	TN contractor run searches around vendors (or TN cities as fallback) (§6)	DONE
Add a TN license source (Nashville open data → TDCI)	Used automatically; Dashboard › Statewide TN license verify toggle for the optional extra check (§3)	DONE
Treat license as validation/enrichment, not discovery	Discovery still = Google + BBB + keywords; license only enriches	DONE
Filter license by trade classification, not the word “drywall”	Handled inside the license source logic	DONE

Workstream C — Vendor / distributor scraper (new)

Requirement (from the brief)	Where to see it in the dashboard	
Build vendor scraper as a separate mode	Dashboard › Run scope › Vendors + its own Vendors page (§7, §8)	DONE
Anchor to the prioritized cities, 20-mile vendor radius	Vendor run circles TN city centers at the vendor radius (§7)	DONE
Vendor alias / subsidiary map (GMS, L&W/ABC, FBM...)	Vendor Aliases page; rolled-up name shown as canonical_network (§8, §9)	DONE
Capture independents by drywall/gypsum category	Vendor relevance filter keeps real distributors, drops contractors (§7)	DONE
Output prioritizes distributor name + branch location	Vendors table + per-vendor detail box (§7, §10)	DONE

Workstream D — Lumber exclusion (3 layers)

Requirement (from the brief)	Where to see it in the dashboard	
Layer 1 — exclude by source category (Google/BBB)	Applied in the Classify stage of every run; removed rows keep an excluded_reason you can read in the detail box, Pipeline Stages and the “include excluded” export (§9, §10, §11)	DONE
Layer 2 — editable negative-keyword list		DONE
Layer 3 — name-pattern matching (clean category, lumber name)		DONE
Apply to both contractor and vendor pipelines	Same filter runs on both lists	DONE
Flag, don't hard-delete — keep an excluded_reason	Rows are tagged and set aside, never deleted (§11)	DONE

Care taken: “plumber” is **not** treated as lumber — the filter matches whole words, not letters.

Workstream E — Data separation & access layer

Requirement (from the brief)	Where to see it in the dashboard	
Stage the pipeline; keep raw data immutable	Pipeline Stages page — each stage's snapshot, viewable & downloadable (§4, §11)	DONE
Separate by tags, not copies (record_type, state, city, tier, zip, source, status flags)	Tag badges on every record + the Tags section of the detail box (§9, §10)	DONE

Deduplicate via entity resolution (canonical_entity_id)	Shown as entity ID in the detail box; links a record to its sources (§9, §10)	DONE
Provenance / auditability — which source produced what	Sources (raw layer) section of the detail box (§10)	DONE
Config & reference data kept out of results tables	Separate editable pages: Cities , Vendor Aliases , Keywords (§8)	DONE
Deliverables are derived — keep data wide, narrow on export	Export CSV on Contractors/Vendors drops lumber/out-of-territory by default; “include excluded” keeps everything (§11)	DONE
Capture everything (wide), let export do the narrowing	Contractors, Vendors and every Pipeline Stage export carry the full column set (all tags)	DONE

HOW THE DATA FLOWS, END TO END

You pick what/where (Dashboard) → the run moves through the **5 stages (Pipeline Stages)** → each business is tagged and deduped into one **canonical record** with its **raw sources** kept underneath → it lands on the **Contractors** or **Vendors** page → you **Export** a clean, wide CSV. Nothing is thrown away; the export just hides what you don't sell to.

13 · Quick answers

"I ran a Tennessee contractor scrape and got vendors / odd results."

Check the Logs. If it used the **Florida-style fallback**, no vendors existed yet. Run a **Tennessee Vendor** scrape first, then run contractors again so the vendor-geography logic can work.

"Why is Florida disabled?"

You're in **Vendors** mode. Vendors are Tennessee-only. Switch to **Contractors** to enable Florida.

"My Florida results look the same as always — is that right?"

Yes. Florida was deliberately left untouched. All the new behaviour is Tennessee/Vendor only.

"What's the difference between Tier 1 and Tier 2?"

Tier is how important a city/market is — Tier 1 is top priority. It's just a label to help you focus; it doesn't remove anything.

"Where did this email/phone come from?"

Click the record → open the **Sources** section. Each source row shows which provider supplied which data.

"How do I get a clean list to hand to sales?"

Open **Contractors** (or **Vendors**), filter if you like, and click **Export CSV**. Lumber and out-of-territory rows are removed automatically.

This guide describes the dashboard as an end user sees it. Engineering details (data sources, internal IDs, license providers) are intentionally kept light here — the goal is to help you run scrapes and read results with confidence.